

**2nd Fantastic Futures Conference**  
Pre-Conference Workshop: Cloud Services Primer  
2019-12-03, Zac Painter, Stanford University

**Lesson Material Credits:**  
The Carpentries (Data Carpentry for Genomics)  
Various Lectures from CS349D at Stanford  
Discussions with IT and HPC staff at Stanford

**What is Cloud Computing?**

Essentially, computing with large systems without direct user management  
You define your parameters and requirements, and the system takes it from there

**Have we done this before?**

Sort of! Mainframe computing used to be standard before personal computers were mainstream  
Despite how powerful personal computers are now, our modern computing needs can outstrip them  
For large computing jobs now, you probably need some kind of extra power than your local machine

**What are my remote computing choices?**

Local Machines/Clusters (e.g., Farmshare/Sherlock at Stanford, MGHPCC in Massachusetts)  
Open Research Compute Networks (e.g., XSEDE, Open Science Grid)  
Commercial Cloud Providers (e.g., Amazon Web Services, Google Cloud, Microsoft Azure)

**What can I get from a cloud? Some or all of the following stack...**

**Software as a Service (SaaS)**  
Centrally-hosted, Subscription-based licensing model (Communication Services, GitHub, etc.)

**Platform as a Service (PaaS)**  
Allow you to work with applications without other barriers (Web Applications and Servers, Databases, etc.)

**Infrastructure as a Service (IaaS)**  
APIs that allow customers to access physical resources (Virtual Machines, Storage, Networks, etc.)

**Cloud/Cluster Architecture**  
Servers with CPU/GPU, RAM, Disks live on...  
Racks of about 40-80, which are combined into...  
Clusters of those for a data center

**Data Lakes**  
Repositories where raw data is stored and manipulated  
More important as compute is moving to where data is

**Other Cloud Differentials**  
  
Edge Computing brings high traffic compute locally  
Fog Computing works locally, using Internet to route

Public Clouds open to anyone to rent  
Private Clouds only for members of specific group

**General Commercial Cloud Pricing Schema**  
Compute pricing usually pay-as-you-go (especially IaaS)  
No minimum/up-front fee  
Most tenants not at 100% use – can be oversold?  
Economies of scale for purchase/power/manage machines  
Not all machines will be the same, some old – some new

**TCO = capital (CapEx) + operational (OpEx) expenses**

**Operators perspective** (Hardware dominates the cost)  
CapEx: building, generators, A/C, compute/storage/net HW  
Including spares, amortized over 3 – 15 years  
OpEx: electricity (5-7c/KWh), repairs, people, WAN, insurance, ...

**Users perspective**  
CapEx: cost of long term leases on HW and services  
OpeEx: pay per use cost on HW and services, people

**The \$64,000 Question... Can you build it cheaper?**

The Answer? Possibly, depending on your scale and scope  
GPU time is a limiting factor; those resources tend to be expensive  
Need to factor in some cost for local maintenance and support

**HPC Clusters**  
Often the cheapest solution if you have it  
Usually fast local help and support options  
Generally fewer privacy/security concerns than public cloud

Most put some kind of limits on what you can do; less flexible  
Storage, Power you can use at once, or Process time often limited  
Most are shared, so you need to use a scheduler (e.g., Slurm)

**Open Researcher Clouds**

Often supported by national research directives (e.g., NSF)  
Usually free or low-cost for researchers with those grants

**Extreme Science and Engineering Discovery Environment (XSEDE)**  
Works similar to HPC Clusters: you need to run schedulers  
You need to bid for time and allocations like a grant  
Upshot is that you can bid for almost anything you need

**Open Science Grid (OSG)**  
Runs off of shared systems at partner institutions; no scheduling  
More limited in overall ability, but a strong choice for some tasks

**Commercial Comparison**  
  
**Overall:**  
Compute prices roughly the same across all three majors  
Service offerings (e.g., computing, storage) broadly similar  
Very hard to support infrastructure for multiple vendors  
Transfer from one platform to another is slow and/or difficult  
You pay for renting the computer/whatever, whether you use it or not

**Amazon Web Services:**  
Probably largest provider at Stanford  
Probably the provider who has invested the most in education/research

**Google Cloud Platform:**  
First to sign at Stanford  
Tends to be more business or third-party service driven  
Probably easiest to use with Kubernetes?

**Microsoft Azure:**  
Generally business/enterprise driven  
Usually slightly more expensive instances than the other two  
If you are Windows driven, this might be cheaper/easier

**Useful Unix commands for you to know**

ssh user\_name@address : How you log into most environments  
whoami : Shows what user and computer you are  
pwd : Shows you where you currently are  
ls : Lists the contents of a directory

df -h : Shows the space on the hard drive  
top : Checks stats of the current system  
ps : Shows all processes in current system

tmux ... : Allows you to keep running jobs when disconnected  
screen ... : An older alternative to tmux on some machines

sudo apt upgrade : Common command to get things ready to install  
sudo apt install ... : Common command to install programs

**CAP Theorem**

In distributed systems, choose 2 out of 3\*  
No network system safe from partition failure  
Choice is really between C & A in the event of partition-failure  
Partitions less common now than before, but still a question  
Serverless computing is increasingly popular  
BASE and ACID are similar ideas

**Consistency**

Every read returns data from most recent write  
Failure can be from bad updates or multiple locations

**Availability**

Every request executes & receives a (non-error) response  
High vs Perfect is generally the major choice

**Partition-tolerance**

The system continues to function when network partitions occur  
If you have partitions, you can't have perfect and consistent

**Containers and Orchestration**

PaaS Services that allow you to replace VMs with portable “containers”  
Allows for some transfer of items from one user/cluster to another  
If you need to move from one environment to another, this is how to do it

Orchestration gathers containers and deploys them automatically  
Docker, Mesos, and Kubernetes are the three largest current players

**Security & Privacy**

Theoretically, service provider can access anything uploaded at any time  
Encrypting data is a solution, but adds extra work

Terms of service sometimes silent on question of “who owns the data?”  
What happens if hardware collapses, or the provider shuts down?  
Clouds can have insecure APIs or Infrastructure, creating risk